

Author Response: CHORD Hallucination Mitigation

We thank the reviewers for identifying that aggregate benchmark scores alone do not isolate mechanism, detector attribution, and deployment cost. We therefore add targeted diagnostics and narrow the claim: CHORD is a detector-assisted, base-MLLM-training-free admission-time verifier for object-grounded hallucination mitigation; decoder-to-vision attention is used as an operational scoring signal, not as a causal explanation.

Concern	Reviewer need	Response / evidence to add	Boundary in final paper
Future mechanism	M8du asks whether Future actually changes admitted tokens and whether those changes help.	We add a paired Full-vs-P+C admission diagnostic: rollout-induced token changes, corrected unsupported mentions, and harmful changes. Exact flip/correctness counts are included only for completed paired logs; otherwise the rebuttal states the diagnostic protocol without expected numbers.	Future is a quality-oriented verifier, not a claim that every token needs rollout. If paired logs are unavailable, we do not report expected percentages, counts, confidence intervals, or p-values.
Detector attribution	KrEs/M8du ask whether Grounding DINO alone explains the gains.	We isolate detector contribution with fixed same-anchor, uniform/random-anchor, no-Current, P+C, and Full controls. If control logs are complete, the compact attribution row is reported; otherwise we state the control design and limitation.	CHORD is detector-assisted. We do not claim detector independence or second-proposer robustness unless those runs are measured.
Cost/defaults and baselines	jjVG/KrEs/ve3y ask about overhead, recent baselines, and why $k = 5, m = 3$.	We separate proposal cost from decode ITL. P+C is the practical/default regime; Full is quality-oriented/offline. $k = 5, m = 3$ is the quality setting, while P+C or $k = 5, m = 2$ is practical under latency constraints. ONLY, VHD/VHR, and HALC use matched backbone, split, prompt family, parser, seeds, and validation budget.	We report numeric total cost, VRAM, batch behavior, and recent-baseline wins only when measured with provenance; otherwise they remain camera-ready additions, not rebuttal evidence.

We will redraw Fig. 2 as sequential Past/Current/Future stages. The camera-ready version will remove detector-independent, broad relation/composition, or causal-attention claims unless supported by measured evidence. All exact numbers in the official page must be measured or removed; expected target values from the internal master are not used as factual rebuttal evidence.