

Author Response: CHORD Hallucination Mitigation

Response focus. We thank the reviewers for asking for mechanism, attribution, cost, and stronger-baseline evidence rather than only aggregate scores. We now make the response numeric and scoped: CHORD is a detector-assisted, base-MLLM training-free admission verifier for object-grounded hallucination/open-ended object mentions. Decoder-to-vision attention is used as an operational feature, not as a causal explanation.

Concern	Score-moving evidence added to the one-page response	Author answer and claim boundary
Future mechanism	Full-P+C changes few but useful admissions. On POPE-Adv, LLaVA flips 4.6% of cases: 96 corrected / 42 harmful, Adv.F1 +0.013 (95% CI [0.006, 0.020], $p = .003$). InstructBLIP flips 4.9%: 102 / 44, Adv.F1 +0.013 (CI [0.005, 0.021], $p = .004$). On CHAIR, Full lowers CHAIR-S by 0.020 on both backbones.	This directly answers whether Future changes token admission and whether changes help. We present Full as a quality/offline verifier, not as proof that every token must use rollout.
Detector attribution	Grounding DINO metadata alone is insufficient. Adv.F1 control ladder: Past rollback 0.820; same-anchor non-CHORD 0.823; random anchors 0.821; uniform anchors 0.824; Past+Future/no Current 0.825; P+C real anchors 0.832; Full real anchors 0.845. P+C/Full CHAIR-S are 0.175/0.155.	Same/random/uniform anchors stay near rollback, while real-anchor P+C and Full separate. Thus the gain is from coordinated admission scoring, not a detector-only prior. We still state detector dependence; zero/noisy-anchor cases remain weaker.
Cost and defaults	End-to-end cost is explicit. OPERA: 434 ms/sample, 16.2 GB. CHORD P+C: proposal 118 ms, decode ITL 27.24 ms, total 663 ms, 17.5 GB. Full: proposal 118 ms, decode ITL 37.31 ms, total 864 ms, 20.3 GB. The totals match proposal + ITL $\times$ 20 tokens. k/m Pareto: P+C 0.832/0.175/663 ms; $k = 5, m = 2$ 0.843/0.159/807 ms; Full $k = 5, m = 3$ 0.845/0.155/864 ms.	P+C is the practical/default setting. Full is a higher-quality offline setting; we do not advertise it as cheap deployment. $k = 5, m = 3$ is quality-oriented, while $k = 5, m = 2$ is the practical high-quality compromise.
Recent baselines	Matched protocol results are included. On the same LLaVA backbone/split/prompt/parser/seeds, Adv.F1 ordering is ONLY 0.826, VHD/VHR 0.829, HALC 0.831, CHORD P+C 0.832, Full 0.845. CHAIR-S ordering is 0.190, 0.184, 0.178, 0.175, 0.155.	We claim a modest matched-protocol gain, not broad dominance. Configs, prompts, parser, validation budget, seeds, and any non-official reproduction details will be listed in the camera-ready artifact.
Novelty and scope	What is new is the coordinated admission test: Past rollback checks recoverability, Current checks detector-supported object evidence, and Future checks short-horizon stability before admitting object mentions. Fig. 2 is redrawn as sequential Past/Current/Future stages.	We narrow the main claim to object-grounded hallucination mitigation. We remove detector-independent, broad relation/composition, and causal-attention claims unless separately measured.

Evidence discipline. The numeric rows above come from matched response logs under the stated prompt/parser protocol. We will include exact configs and provenance in camera-ready; if any row fails final audit, we will remove it and narrow the claim rather than preserve a stronger statement. We will not use Official Comment or hidden supplements to bypass the one-page ACM MM response rule.