

CHORD ACM MM Rebuttal Draft with Reviewer-Audited Expected Tables

Anonymous Submission 8826

Internal draft, 2026-05-29

Critical boundary. This PDF is an internal rebuttal draft and experiment target. Tables marked *expected*, *target*, or *expected placeholder* are not measured rebuttal results. Before submission, every such entry must be replaced by measured SSH experiment outputs or the corresponding claim must be weakened.

Format boundary. The public ACM MM 2026 page states that rebuttals are submitted in OpenReview, must remain anonymous, and cannot include external links. OpenReview’s default rebuttal form is Markdown text with max length 2500 characters; however, the logged-in ACM MM 2026 form controls the final length and whether a PDF is accepted. This PDF should therefore be treated as a planning and compression artifact, not as proof that a PDF upload is allowed.

1. Rebuttal Goal

The reviews are favorable on topic fit but skeptical about attribution. The rebuttal should not claim that every ingredient is new in isolation. The strongest defensible thesis is:

CHORD is a detector-assisted, training-free-for-the-base-MLLM admission-time verifier. Past+Current is the latency-oriented regime; Full CHORD is the quality-oriented regime, especially when open-ended continuation stability matters.

2. Reviewer Concern Map

Reviewer	Core blocker	Evidence that can move the score	Rebuttal stance
jjVG	Cost, Figure 2, k/m, missing ONLY/VHD	End-to-end cost, k/m Pareto sweep, related-work axes	Acknowledge and commit to concise revisions.
KrEs	Novelty and Grounding DINO attribution	Real/no/random anchors, same-anchor non-CHORD, cost/VRAM	Show CHORD scoring adds value beyond detector metadata.
yx8u	Incremental novelty, detector dependence, attention reliability	Claim calibration plus detector/failure strata	Use “operational signal”, not causal explanation.
ve3y	Moderate novelty and overhead	Two-regime cost-quality frontier	Frame Full as quality mode, not cheap default.
M8du	Future mechanism and flip correctness	Full vs P+C corrected/harmful flips, CHAIR effect	Answer mechanism first; do not rely on aggregate scores only.

3. Submitted Numeric Anchors

These are already in the submitted paper and can anchor the expected ranges. They are not new rebuttal diagnostics.

Model	Regime	Adv. F1	CHAIR-S	CHAIR-I	MMB	ITL
LLaVA-1.5	Greedy	0.8036	0.2291	0.2046	68.63	19.73
LLaVA-1.5	P+C	0.8318	0.1745	0.1852	69.15	27.24
LLaVA-1.5	Full	0.8453	0.1548	0.1754	69.78	37.31
InstructBLIP	Greedy	0.8327	0.2140	0.3380	69.34	16.47
InstructBLIP	P+C	0.8517	0.1543	0.2946	70.08	24.51
InstructBLIP	Full	0.8651	0.1347	0.2795	70.82	35.86

4. Core Expected Tables

Table A: Future Mechanism

Run	Expected band	PASS threshold	WARN threshold	FAIL threshold
POPE-Adv: Full vs P+C	Answer flip rate 1–8%; modest net correction, not a large perturbation story.	Corrected flips exceed harmful flips by at least 5 per 1000 samples and Adv. F1 improves by at least 0.006 on at least one backbone without a drop on the other.	Small positive margin or Adv. F1 gain 0.002–0.006.	Corrected flips do not exceed harmful flips, or the larger run repeats the 64-sample warning pattern.
CHAIR: Full vs P+C	Submitted CHAIR-S gain is about 0.020, so a matched expected band is 0.012–0.025.	CHAIR-S improves by at least 0.012 and CHAIR-I does not worsen by more than 0.005.	CHAIR-S improves 0.005–0.012 or only one backbone is clean.	CHAIR-S improves below 0.005 or worsens.
Past+Future P+C vs Full	Full should beat one-missing variants on open-ended quality if current and future are complementary.	Full best or tied best on CHAIR-S with no major POPE/retention harm.	Full helps only on CHAIR.	P+C or Past+Future dominates Full on both POPE and CHAIR.

Reviewer-safe wording if Table A passes: Future is a bounded continuation-stability signal, not a universal yes/no correction. If Table A fails, the rebuttal must lead with Past+Current and remove Future as a main mechanism claim.

Table B: Detector Attribution

Control	What it isolates	Expected target	Stop rule
Real vs uniform anchors	Whether query-conditioned spatial weighting matters.	Real anchors beat uniform by at least 0.006 Adv. F1 or 0.012 CHAIR-S.	If uniform matches real, do not claim object-resonant grounding.
Real vs random matched anchors	Whether arbitrary visual weighting explains gains.	Real anchors beat random; random has worse false positives or CHAIR.	If random equals or beats real, rewrite as detector-assisted attention regularization at most.
Full vs Past+Future	Whether current grounding still matters when Future is present.	Full improves CHAIR-S by at least 0.010 or Adv. F1 by 0.006.	If Past+Future matches Full, the coordination claim is weak.
Same-anchor non-CHORD	Whether detector metadata alone explains gains.	P+C/Full exceed same-anchor control by at least 0.006 Adv. F1 or 0.012 CHAIR-S.	If same-anchor control matches CHORD, concede detector attribution.
Failure strata	Robustness under zero/noisy/missed anchors.	Relevant-anchor strata improve; zero/noisy anchors are bounded and acknowledged.	If noisy/zero-anchor cases drive gains, remove robustness claim.

Table C: End-to-End Cost

Method	Required measured fields	Expected interpretation	Red flag
Greedy OPERA	Decode ITL, total ms/sample, generated tokens, peak VRAM.	Baseline and rollback-only cost floor.	Missing generated-token normalization.
Past+Current	Proposal ms/img, decode ITL, total latency, peak VRAM, batch size.	Practical regime if transparent and around or below 1.8–2.2x Greedy total latency depending on answer length.	Detector time hidden or P+C exceeds 2.2x on CHAIR.
Full	Same fields as P+C.	Quality mode; acceptable only with clear CHAIR/mechanism gain.	Full exceeds 3.2x Greedy total without strong CHAIR gain.
Batch behavior	Batch 1 required; batch 2/4 optional if memory allows.	Shows deployment boundary honestly.	Peak VRAM omitted while making practical claims.

Table D: Hyperparameter Robustness

Sweep	Expected result	Rebuttal rule
$k \in \{3, 5, 10\}$, $m \in \{1, 2, 3, 4\}$	$k = 5, m = 3$ should be Pareto-near, not necessarily strictly best. Larger k or m should show diminishing returns.	Defend default only if no other setting improves Adv. F1 by 0.006 or CHAIR-S by 0.010 while reducing total latency by at least 10%.
$m = 0$ / P+C reference	Low-cost baseline for showing Future contribution.	If P+C dominates Full, make P+C the default claim.

Table E: Related-Work Positioning

Work	Safe positioning	Numeric comparison rule
OPERA, VCD, DoLa	Already submitted direct decode-time baselines.	Existing submitted numbers can be cited.
HALC	Plug-and-play hallucination mitigation baseline; closest if matched implementation exists.	Report only if same model, prompts, parser, and splits are reproducible.
ONLY	Reviewer-named recent training-free method; must be added to related work.	No number unless code/settings are verified.
VHD/VHR	Reviewer-named vision/head-divergence family; must be added to related work after exact identity verification.	No memory-based citation or unmatched number.

5. OpenReview-Compressed Rebuttal Skeleton

The final OpenReview answer should be rebuilt from measured results. Until then, placeholders must remain visibly marked:

We thank the reviewers. The shared concern is whether CHORD’s gains come from the coordinated admission rule rather than detector metadata or extra compute. We therefore prepared targeted diagnostics. First, Full vs P+C should change only a small fraction of POPE-Adv answers *Expected placeholder: 1 to 8%*, with corrected flips exceeding harmful flips *Expected placeholder: at least +5 corrected per 1000 samples* and a clearer CHAIR-S improvement *Expected placeholder: 0.012 to 0.025 absolute reduction*. If measured, this would support the revised claim that Future is a bounded continuation-stability signal, not a universal yes/no correction. Second, detector controls should show real anchors outperform random/uniform anchors *Expected placeholder: at least +0.006 Adv. F1 or -0.012 CHAIR-S* and a same-anchor non-CHORD control remaining below CHORD; otherwise we will narrow the detector-attribution claim. Third, we will report proposal time, decode ITL, total latency, generated length, peak VRAM, and batch size; P+C is the latency-oriented regime and Full is the quality-oriented regime. Fourth, a k/m sweep should show $k = 5, m = 3$ is Pareto-near *Expected placeholder: no alternative improves Adv. F1 by 0.006 or CHAIR-S by 0.010 while reducing total latency by at least 10%*. We will revise wording to "detector-assisted and training-free for the base MLLM", add ONLY/VHD/HALC positioning, redraw Figure 2, and treat attention as an operational scoring feature rather than a causal explanation.

6. Final Submission Checklist

- Replace every expected placeholder and every expected band with measured JSON/JSONL-backed values.
- If any P0 stop rule fails, narrow the claim instead of arguing around it.
- Verify the logged-in ACMMM 2026 OpenReview form before deciding whether to submit text only or attach a PDF.
- Keep the response anonymous and include no external links.
- Do not cite the 64-sample POPE random diagnostic as positive evidence for Full CHORD.